

Tensorflow ConvNet

MNIST Classification
CAT vs DOG Classification

최준명

Pusan National University

Vision & Intelligent system Lab

<https://jimmylifestudy.blogspot.kr/>

jmchoi@pusan.ac.kr

Tensorflow | 지난 시간 복습

- 먼저 **Graph**를 Build하고 **Session**을 통해 실행한다.
- 변수는 `tf.Variable()` 보단 `tf.get_variable()`
- 입력은 `tf.placeholder()` with `feed_dict` 나 `tf.data()`를 쓰자.
- 모델 building은 **Layer API** 를 쓰면 간편하다. (단점: debugging이 어려움)
- 간단한 Linear Regression 모델
- 간단한 Logistic Regression 모델

Tensorflow | 오늘 목표

- Convolutional layer을 이용한 MNIST 분석
- 나만의 데이터를 이용하여 ConNet 구성

Convolutional layer in Tensorflow

Tensorflow | Convolutional layer

MNIST Dataset

$X = 28 \times 28$ array (1d size of 784)

$Y = \text{Digit value}$

MNIST.train: 55,000 examples

MNIST.validation: 5,000 examples

MNIST.test: 10,000 examples

Inference: $Y_{\text{predicted}} = \text{softmax}(X * w + b)$

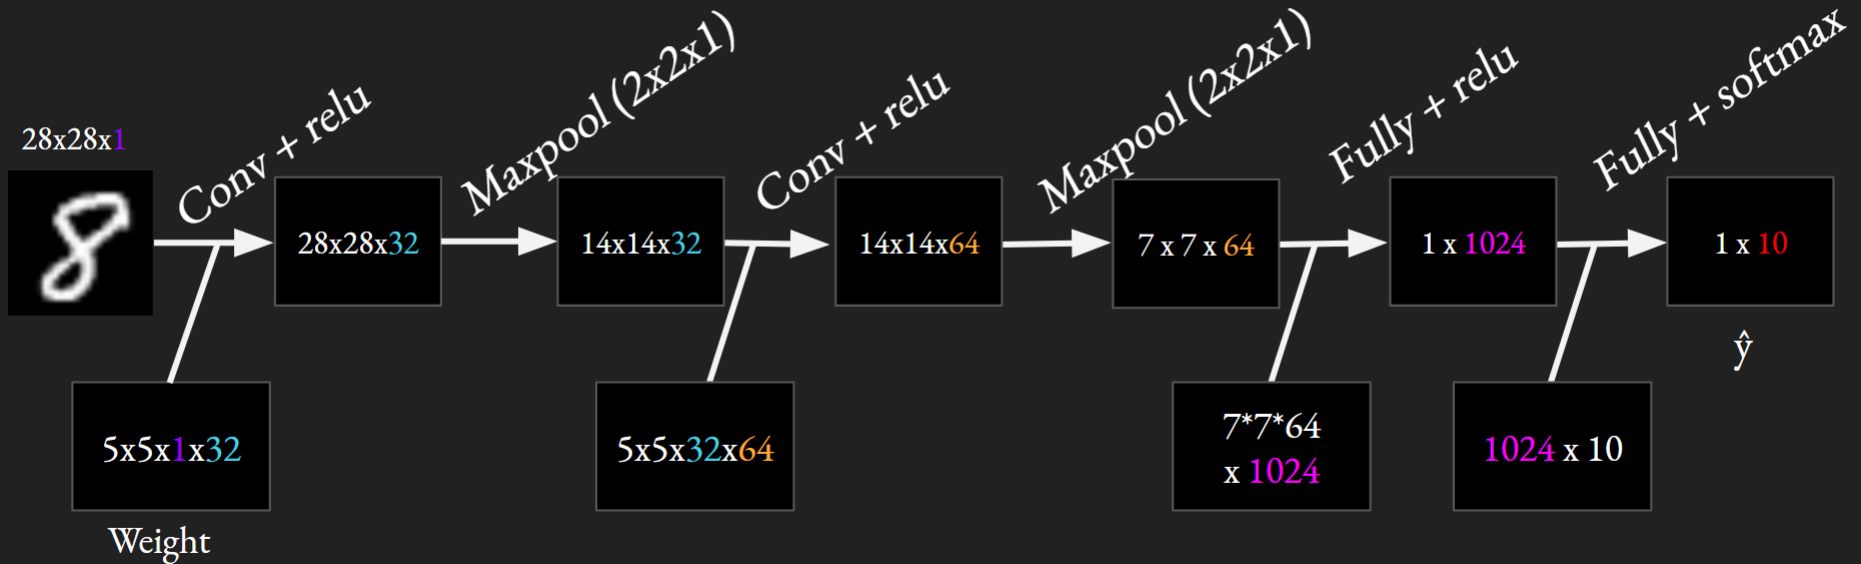
Cross entropy loss:
 $-\log(Y_{\text{predicted}})$



출처: https://docs.google.com/presentation/d/1ImcQVNAJrL8x3lq0VB1mVaka1r6pOlb-TMVTX5RuFc/edit#slide=id.g1c166da651_0_5

Tensorflow | Convolutional layer

Model



출처: https://docs.google.com/presentation/d/17VTArfQVtapBqfYecyvp3Kp9HKy8Pw2WI12acYME2n/edit#slide=id.g1c60f09bdb_0_336

Tensorflow | Convolutional layer

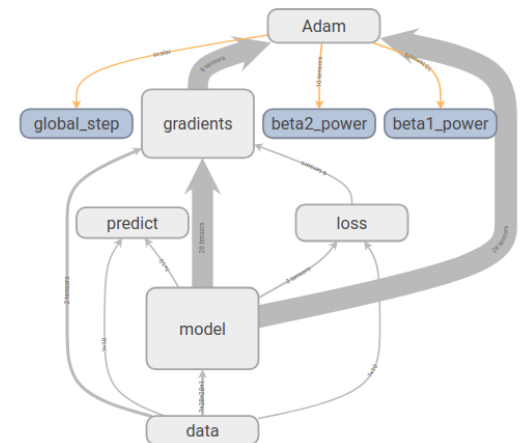
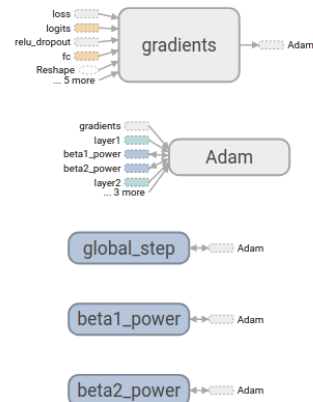
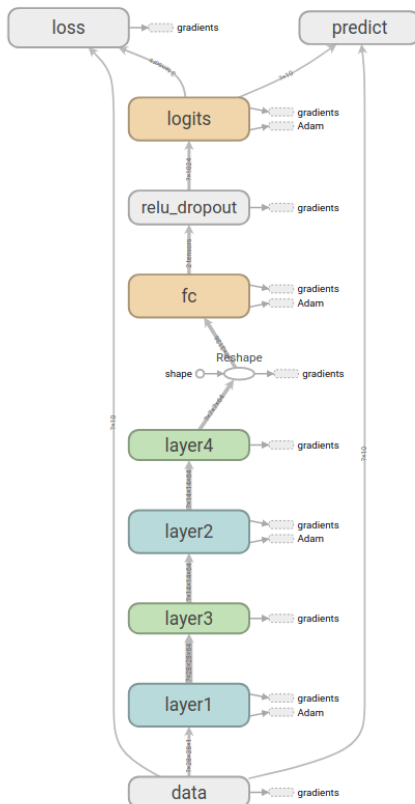
```
def conv2d(inputs, nb_filter, filter_size, strides=1, padding='same',
          activation=tf.nn.relu, bias=True, weights_init=variance_scaling_initializer,
          bias_init=zeros_initializer, trainable=True, scope='Conv2d'):
    with tf.variable_scope(scope, values=[inputs], reuse=tf.AUTO_REUSE) as scope:
        W_init = weights_init
        bias_init = bias_init
        x = tf.layers.conv2d(inputs, filters=nb_filter, kernel_size=filter_size, strides=[strides, strides], padding=padding,
                             kernel_initializer=W_init, bias_initializer=bias_init, activation=activation, trainable=trainable)
    return x

def maxpool(inputs, filter_size=2, strides=1, padding='same', scope='pool'):
    with tf.variable_scope(scope, reuse=tf.AUTO_REUSE) as scope:
        x = tf.layers.max_pooling2d(inputs, pool_size=[filter_size, filter_size], strides=[strides, strides], padding=padding)
    return x

def fully_connected(inputs, nb_filter, activation=tf.nn.relu, bias=True, weights_init=variance_scaling_initializer,
                    bias_init=zeros_initializer, trainable=True, scope='FC'):
    if len(inputs.get_shape().as_list()) > 2:
        n_inputs = int(np.prod(inputs.get_shape().as_list()[1:]))
        inputs = tf.reshape(inputs, [-1, n_inputs])
    with tf.variable_scope(scope, reuse=tf.AUTO_REUSE) as scope:
        x = tf.layers.dense(inputs, units=nb_filter, activation=activation, use_bias=bias, kernel_initializer=weights_init,
                             bias_initializer=bias_init, trainable=trainable)
    return x
```

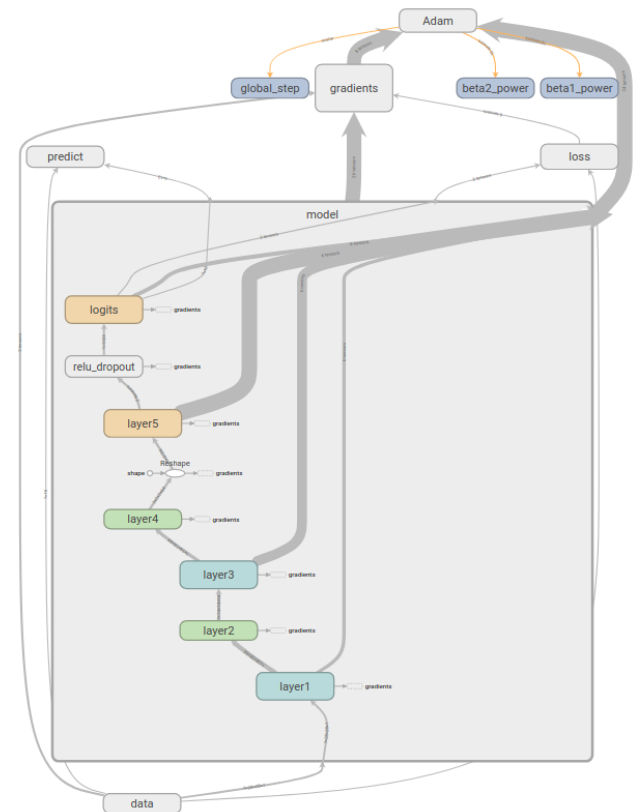
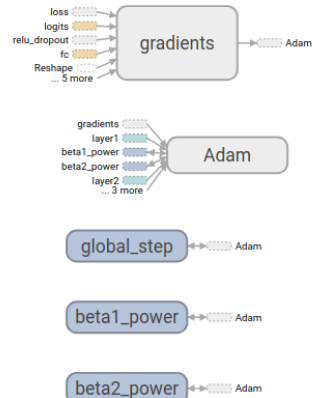
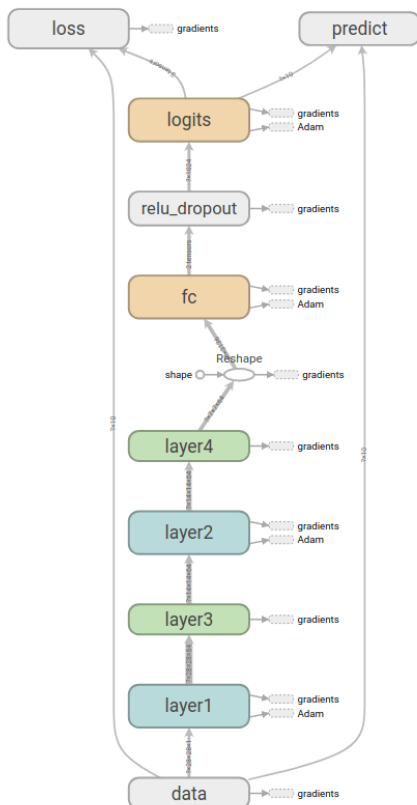
Tensorflow | Convolutional layer

```
def inference(self):  
    with tf.variable_scope('model') as scope:  
        conv1 = conv2d(inputs=self.img, nb_filter=32, filter_size=5, trainable=self.training, scope='layer1')  
        pool1 = maxpool(conv1, 2, 2, 'same', scope='layer2')  
        conv2 = conv2d(inputs=pool1, nb_filter=64, filter_size=5, trainable=self.training, scope='layer3')  
        pool2 = maxpool(conv2, 2, 2, 'same', scope='layer4')  
        fc = fully_connected(pool2, 1024, trainable=self.training, scope='layer5')  
        self.logits = fully_connected(fc, self.n_classes, trainable=self.training, scope='logits')
```



Tensorflow | Convolutional layer

```
def inference(self):  
    with tf.variable_scope('model') as scope:  
        conv1 = conv2d(inputs=self.img, nb_filter=32, filter_size=5, trainable=self.training, scope='layer1')  
        pool1 = maxpool(conv1, 2, 2, 'same', scope='layer2')  
        conv2 = conv2d(inputs=pool1, nb_filter=64, filter_size=5, trainable=self.training, scope='layer3')  
        pool2 = maxpool(conv2, 2, 2, 'same', scope='layer4')  
        fc = fully_connected(pool2, 1024, trainable=self.training, scope='layer5')  
        self.logits = fully_connected(fc, self.n_classes, trainable=self.training, scope='logits')
```



Tensorflow | Convolutional layer

```
def loss(self):
    """
    ~~~~~
    define loss function
    use softmax cross entropy with logits as the loss function
    compute mean cross entropy, softmax is applied internally
    """
    #
    with tf.name_scope('loss'):
        entropy = tf.nn.softmax_cross_entropy_with_logits(labels=self.label, logits=self.logits)
        self.loss = tf.reduce_mean(entropy, name='loss')

def optimize(self):
    """
    ~~~~~
    Define training op
    using Adam Gradient Descent to minimize cost
    """
    self.opt = tf.train.AdamOptimizer(self.lr).minimize(self.loss,
                                                         global_step=self.gstep)

def summary(self):
    """
    ~~~~~
    Create summaries to write on TensorBoard
    """
    with tf.name_scope('summaries'):
        tf.summary.scalar('loss', self.loss)
        tf.summary.scalar('accuracy', self.accuracy_holder)
        tf.summary.histogram('histogram loss', self.loss)
        self.summary_op = tf.summary.merge_all()

def build(self):
    """
    ~~~~~
    Build the computation graph
    """
    self.get_data()
    self.inference()
    self.loss()
    self.optimize()
    self.eval()
    self.summary()
```

Tensorflow | Convolutional layer

```
def train_one_epoch(self, sess, saver, writer, epoch, step):
    start_time = time.time()
    self.training = True
    n_batches = int(self.mnist.train.num_examples/self.batch_size)
    total_loss = 0
    for j in range(n_batches):
        self.X_batch, self.Y_batch = self.mnist.train.next_batch(self.batch_size)
        _, l = sess.run([self.opt, self.loss], feed_dict={self.X: self.X_batch, self.label: self.Y_batch})
        step += 1
        total_loss += l
    saver.save(sess, 'checkpoints/convnet_mnist/mnist-convnet', step)
    print('Average loss at epoch {0}: {1} Took: {2} seconds'.format(epoch, total_loss/n_batches, time.time()-start_time))
    return step

def eval_once(self, sess, writer, epoch, step):
    start_time = time.time()
    self.training = False
    total_correct_preds = 0
    n_batches = int(self.mnist.test.num_examples/100)
    for j in range(n_batches):
        self.X_batch, self.Y_batch = self.mnist.test.next_batch(100)
        accuracy_batch = sess.run(self.accuracy, feed_dict={self.X: self.X_batch, self.label: self.Y_batch})
        total_correct_preds += accuracy_batch
    summaries = sess.run(self.summary_op, feed_dict={self.X: self.X_batch, self.label: self.Y_batch,
                                                    self.accuracy_holder: total_correct_preds/n_batches})
    writer.add_summary(summaries, global_step=step)
    print('Accuracy at epoch {0}: {1} Took: {2} seconds'.format(epoch, total_correct_preds/n_batches,
                                                                time.time()-start_time))
```

Tensorflow | Convolutional layer

```
def train(self, n_epochs):  
    """  
    The train function alternates between training one epoch and evaluating  
    """  
    utils.safe_mkdir('checkpoints')  
    utils.safe_mkdir('checkpoints/convnet_mnist')  
    writer = tf.summary.FileWriter('./graphs/convnet', tf.get_default_graph())  
  
    with tf.Session() as sess:  
        sess.run(tf.global_variables_initializer())  
        saver = tf.train.Saver()  
        ckpt = tf.train.get_checkpoint_state(os.path.dirname('checkpoints/convnet_mnist/checkpoint'))  
        if ckpt and ckpt.model_checkpoint_path:  
            saver.restore(sess, ckpt.model_checkpoint_path)  
  
        step = self.gstep.eval()  
  
        for epoch in range(n_epochs):  
            step = self.train_one_epoch(sess, saver, writer, epoch, step)  
            self.eval_once(sess, writer, epoch, step)  
    writer.close()
```


Tensorflow | Custom dataset

Kaggle dogs vs cats (download [here](#))

X = 크기가 다양함

Y = 0 or 1

dog.1.jpg

dog.2.jpg

....

cat.1.jpg

cat.2.jpg

...



Tensorflow | Custom dataset

```
def get_files(data_path, datafolder):  
    cats = []  
    label_cats = []  
    dogs = []  
    label_dogs = []  
    for file in datafolder:  
        name = file.split('.')  
        if name[0]=='cat':  
            cats.append(os.path.join(data_path, file))  
            label_cats.append(0)  
        else:  
            dogs.append(os.path.join(data_path, file))  
            label_dogs.append(1)  
  
    image_list = np.hstack((cats, dogs))  
    label_list = np.hstack((label_cats, label_dogs))  
  
    shuffle = list(zip(image_list, label_list))  
    random.shuffle(shuffle)  
    image_list, label_list = zip(*shuffle)  
    return image_list, label_list
```

Tensorflow | Custom dataset

```
def get_batch(image, label, image_W, image_H, batch_size, capacity):  
    image = tf.cast(image, tf.string)  
    label = tf.cast(label, tf.int32)  
  
    # make an input queue  
    input_queue = tf.train.slice_input_producer([image, label])  
    label = input_queue[1]  
    image_contents = tf.read_file(input_queue[0])  
    image = tf.image.decode_jpeg(image_contents, channels=3)  
  
    # data argumentation  
    image = tf.image.resize_image_with_crop_or_pad(image, image_W, image_H)  
  
    image = tf.image.per_image_standardization(image)  
  
    image_batch, label_batch = tf.train.batch([image, label],  
                                              batch_size=batch_size,  
                                              num_threads=64,  
                                              capacity=capacity)  
    label_batch = tf.one_hot(label_batch, depth=2)  
    image_batch = tf.cast(image_batch, tf.float32)  
  
    return image_batch, label_batch
```

Tensorflow | Custom dataset

Data augmentation (https://www.tensorflow.org/api_docs/python/tf/image)

`adjust_brightness(...)` : Adjust the brightness of RGB or Grayscale images.

`adjust_contrast(...)` : Adjust contrast of RGB or grayscale images.

`adjust_gamma(...)` : Perform

`adjust_hue(...)` : Adjust hue

`adjust_saturation(...)` : A

`central_crop(...)` : Crop the

`convert_image_dtype(...)` :

`crop_and_resize(...)` : Extr

`crop_to_bounding_box(...)`

`decode_and_crop_jpeg(...)`

`decode_bmp(...)` : Decode the first frame of a BMP-encoded image to a uint8 tensor.

`decode_gif(...)` : Decode the first frame of a GIF-en

`decode_image(...)` : Convenience function for deco

`decode_jpeg(...)` : Decode

`decode_png(...)` : Decode

`draw_bounding_boxes(...)`

`encode_jpeg(...)` : JPEG-e

`encode_png(...)` : PNG-enc

`extract_glimpse(...)` : Ex

`extract_jpeg_shape(...)`

`flip_left_right(...)`

`flip_up_down(...)` : Fl

`grayscale_to_rgb(...)`

`hsv_to_rgb(...)` : Con

`is_jpeg(...)` : Conven

`non_max_suppression(...)`

`pad_to_bounding_box(...)`

`per_image_standardiz`

`random_brightness(...)`

`random_contrast(...)`

`random_flip_up_down(...)` : Randomly flips an image vertically (ups

`random_hue(...)` : Adjust the hue of an RGB image by a random facto

`random_saturation(...)` : Adjust the saturation of an RGB image by

`resize_area(...)` : Resize images to size using area interpolation

`resize_bicubic(...)` : Resize images to size using bicubic interp

`resize_bilinear(...)` : Resize images to size using bilinear inter

`resize_image_with_crop_or_pad(...)` : Crops and/or pads an imag

`resize_images(...)` : Resize images to size using the specified m

`resize_nearest_neighbor(...)` : Resize images to size using ne

`rgb_to_grayscale(...)` : Converts one or more images from RGB to

`rgb_to_hsv(...)` : Converts one or more images from RGB to HSV.

`rgb_to_yiq(...)` : Converts one or more images from RGB to YIQ.

`rgb_to_yuv(...)` : Converts one or more images from RGB to YUV.

`rot90(...)` : Rotate image(s) counter-clockwise by 90 degrees.

`sample_distorted_bounding_box(...)` : Generate a single randoml

Tensorflow | Convolutional layer

```
with tf.variable_scope('conv1') as scope:
    weights = tf.get_variable('weights',
                              shape=[3, 3, 3, 16],
                              dtype=tf.float32,
                              initializer=tf.truncated_normal_initializer(stddev=0.1, dtype=tf.float32))
    biases = tf.get_variable('biases',
                              shape=[16],
                              dtype=tf.float32,
                              initializer=tf.constant_initializer(0.1))
    conv = tf.nn.conv2d(images, weights, strides=[1, 1, 1, 1], padding='SAME')
    pre_activation = tf.nn.bias_add(conv, biases)
    conv1 = tf.nn.relu(pre_activation, name=scope.name)

# pool1
with tf.variable_scope('pooling1_lrn') as scope:
    pool1 = tf.nn.max_pool(conv1, ksize=[1, 3, 3, 1], strides=[1, 2, 2, 1],
                           padding='SAME', name='pooling1')

# conv2
with tf.variable_scope('conv2') as scope:
    weights = tf.get_variable('weights',
                              shape=[3, 3, 16, 16],
                              dtype=tf.float32,
                              initializer=tf.truncated_normal_initializer(stddev=0.1, dtype=tf.float32))
    biases = tf.get_variable('biases',
                              shape=[16],
                              dtype=tf.float32,
                              initializer=tf.constant_initializer(0.1))
    conv = tf.nn.conv2d(pool1, weights, strides=[1, 1, 1, 1], padding='SAME')
    pre_activation = tf.nn.bias_add(conv, biases)
    conv2 = tf.nn.relu(pre_activation, name='conv2')

# pool2_a
with tf.variable_scope('pooling2_lrn') as scope:
    pool2 = tf.nn.max_pool(conv2, ksize=[1, 3, 3, 1], strides=[1, 1, 1, 1],
                           padding='SAME', name='pooling2')
```

Tensorflow | Convolutional layer

```
# local3
with tf.variable_scope('local3') as scope:
    reshape = tf.reshape(pool2, shape=[batch_size, -1])
    dim = reshape.get_shape()[1].value
    weights = tf.get_variable('weights',
                              shape=[dim, 128],
                              dtype=tf.float32,
                              initializer=tf.truncated_normal_initializer(stddev=0.005, dtype=tf.float32))
    biases = tf.get_variable('biases',
                              shape=[128],
                              dtype=tf.float32,
                              initializer=tf.constant_initializer(0.1))
    local3 = tf.nn.relu(tf.matmul(reshape, weights)+biases, name=scope.name)

# local4
with tf.variable_scope('local4') as scope:
    weights = tf.get_variable('weights',
                              shape=[128, 128],
                              dtype=tf.float32,
                              initializer=tf.truncated_normal_initializer(stddev=0.005, dtype=tf.float32))
    biases = tf.get_variable('biases',
                              shape=[128],
                              dtype=tf.float32,
                              initializer=tf.constant_initializer(0.1))
    local4 = tf.nn.relu(tf.matmul(local3, weights)+biases, name='local4')

# softmax
with tf.variable_scope('softmax_linear') as scope:
    weights = tf.get_variable('softmax_linear',
                              shape=[128, n_classes],
                              dtype=tf.float32,
                              initializer=tf.truncated_normal_initializer(stddev=0.005, dtype=tf.float32))
    biases = tf.get_variable('biases',
                              shape=[n_classes],
                              dtype=tf.float32,
                              initializer=tf.constant_initializer(0.1))
    softmax_linear = tf.add(tf.matmul(local4, weights), biases, name='softmax_linear')
```

Tensorflow | Custom dataset

```
def _run():
    train, train_label = get_files(data_path=data_path, datafolder=datafolder)

    train_batch, train_label_batch = get_batch(train,
                                                train_label,
                                                img_W,
                                                img_H,
                                                batch_size,
                                                batch_size*2)

    train_logits = inference(train_batch, batch_size, n_classes)
    train_loss = losses(train_logits, train_label_batch)
    train_op = training(train_loss, learning_rate)
    train_acc = evaluation(train_logits, train_label_batch)

    summary_op = tf.summary.merge_all()
    sess = tf.Session()
    train_writer = tf.summary.FileWriter(logs_train_dir, sess.graph)
    saver = tf.train.Saver()

    sess.run(tf.global_variables_initializer())
    coord = tf.train.Coordinator()
    threads = tf.train.start_queue_runners(sess=sess, coord=coord)

    try:
        for step in np.arange(10000):
            if coord.should_stop():
                break
            tra_loss, tra_acc = sess.run([train_op, train_loss, train_acc])
```

Tensorflow | 요약

- ConvNet 기반 MNIST 분석
- 나만의 데이터를 이용한 데이터 분석
- 다양한 방식의 코드 분석

To be continued..
